

AWS

3 REPORT

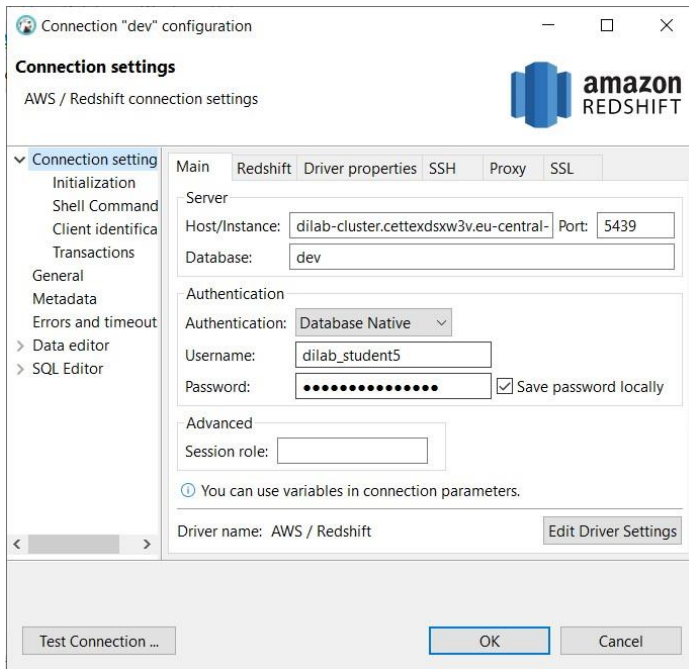
CONTENTS

CREATING THE REPORT	2
TASK 1. LOGIN TO REDSHIFT	2
TASK 2 COPY COMMAND	2
TASK 3	5
CHECK COMPRESSION TYPES FOR NO_COMP	6
COMPARE SIZE OF TABLES WITH DEFAULT COMPRESSION, WITHOUT COMPRESSION AND ANALYZED	7
TASK 4. WORK WITH STORED PROCEDURE	7
TASK 5. WORK WITH OPTIMIZATION OF DISTRIBUTION STYLE AND SORT KEYS.	10
TASK 6. RUN THE STORED PROCEDURE USING OPTIMIZED TABLES AND LOAD DATA TO YOUR REPORT.	11
COPY QUESTION	12
EXTERNAL TABLES	14

Legal Notice: This document contains privileged and/or confidential information and may not be disclosed, distributed or reproduced without the prior written permission of EPAM®.

CREATING THE REPORT

LOGIN TO REDSHIFT



The screenshot shows the 'Connection "dev" configuration' dialog box for Amazon Redshift. The 'Main' tab is selected, showing the 'Server' section with 'Host/Instance' set to 'dilab-cluster.cettexdsw3v.eu-central-1.amazonaws.com' and 'Port' set to '5439'. The 'Database' is 'dev'. In the 'Authentication' section, 'Authentication' is set to 'Database Native', 'Username' is 'dilab_student5', and 'Password' is masked with dots. The 'Save password locally' checkbox is checked. The 'Advanced' section has 'Session role' empty. At the bottom, 'Driver name' is 'AWS / Redshift' and there is an 'Edit Driver Settings' button. The 'Test Connection ...' button is on the left, and 'OK' and 'Cancel' buttons are on the right.

COPY COMMAND

A. Loading needed tables

```
CREATE SCHEMA user_dilab_student5;

-----CREATING CUSTOMERS IN RS FROM S3

CREATE TABLE user_dilab_student5.dim_customers
(
  customer_surr_id integer NOT NULL,
  source_system VARCHAR(50) NOT NULL,
  source_entity VARCHAR(50) NOT NULL,
  customer_src_id VARCHAR(50) NOT NULL,
  customer_first_name VARCHAR(50) NOT NULL,
  customer_last_name VARCHAR(50) NOT NULL,
  customer_gender VARCHAR(50) NOT NULL,
  customer_birthday DATE NOT NULL,
  customer_phone VARCHAR(50) NOT NULL,
  update_dt DATE NOT NULL,
  insert_dt DATE NOT NULL
);
Commit;

ALTER TABLE user_dilab_student5.dim_customers ADD CONSTRAINT dim_customers_pk PRIMARY KEY ( customer_surr_id );

copy user_dilab_student5.dim_customers (CUSTOMER_SURR_ID,SOURCE_SYSTEM,SOURCE_ENTITY,CUSTOMER_SRC_ID,CUSTOMER_F
from 's3://artiom-dolzhenko-aws-bucket/stores/bl dm/dim_customers/dim_customers.scv'
credentials
'aws_iam role=arn:aws:iam::260586643565:role/dilab-redshift-role'
region 'eu-central-1'
delimiter ','
csv
DATEFORMAT AS 'DD-MON-YY'
IGNOREHEADER 1;
Commit;
```

```
select pg_table_def.schemaname, pg_table_def.tablename, pg_table_def.column, pg_table_def.type,
pg_table_def.encoding, pg_table_def.distkey, pg_table_def.sortkey
from pg_table_def
where schemaname='user_dilab_student5'
and tablename='dim_customers';
```

pg_table_def 1 × Results 0 ×

select pg_table_def.schemaname, pg_table_def.tablename, pg_table_def.col Enter a SQL expression to filter results (use Ctrl+Space)

	asc schemaname	asc tablename	asc column	type	asc encoding	distkey	sortkey
1	user_dilab_student5	dim_customers	customer_surr_id	integer	az64	[]	0
2	user_dilab_student5	dim_customers	source_system	character varying(50)	lzo	[]	0
3	user_dilab_student5	dim_customers	source_entity	character varying(50)	lzo	[]	0
4	user_dilab_student5	dim_customers	customer_src_id	character varying(50)	lzo	[]	0
5	user_dilab_student5	dim_customers	customer_first_name	character varying(50)	lzo	[]	0
6	user_dilab_student5	dim_customers	customer_last_name	character varying(50)	lzo	[]	0
7	user_dilab_student5	dim_customers	customer_gender	character varying(50)	lzo	[]	0
8	user_dilab_student5	dim_customers	customer_birthday	date	az64	[]	0
9	user_dilab_student5	dim_customers	customer_phone	character varying(50)	lzo	[]	0
10	user_dilab_student5	dim_customers	update_dt	date	az64	[]	0
11	user_dilab_student5	dim_customers	insert_dt	date	az64	[]	0

```
-----check distribution style-----
select "schema", "table", diststyle
from SVV_TABLE_INFO
where "schema"='user_dilab_student5'
and "table"='dim_customers';
```

```
CREATE TABLE user_dilab_student5.dim_dates (
date id date not null
```

Results 1 × Results 0

select "schema", "table", diststyle from SVV_TABLE_INFO where "schema"='u Enter a SQL expression to filter results (use Ctrl+Space)

	asc schema	asc table	asc diststyle
1	user_dilab_student5	dim_customers	AUTO(EVEN)

B

-----check compression types, distribution keys, sort keys-----

```
select pg_table_def.schemaname, pg_table_def.tablename, pg_table_def.column, pg_table_def.type,
pg_table_def.encoding, pg_table_def.distkey, pg_table_def.sortkey
from pg_table_def where
schemaname='user_dilab_student5' and
tablename='fct_sales';
```

pg_table_def 1 x

select pg_table_def.schemaname, pg_table_def.tablename, pg_table_def.coli Enter a SQL expression to filter results (use Ctrl+Space)

	schemaname	tablename	column	type	encoding	distkey	sortkey
1	user_dilab_student5	fct_sales	source_system	character varying(50)	lzo	[]	0
2	user_dilab_student5	fct_sales	source_entity	character varying(50)	lzo	[]	0
3	user_dilab_student5	fct_sales	product_surr_id	integer	az64	[]	0
4	user_dilab_student5	fct_sales	payment_type_surr_id	integer	az64	[]	0
5	user_dilab_student5	fct_sales	promotion_surr_id	integer	az64	[]	0
6	user_dilab_student5	fct_sales	store_surr_id	integer	az64	[]	0
7	user_dilab_student5	fct_sales	employee_surr_id	integer	az64	[]	0
8	user_dilab_student5	fct_sales	customer_surr_id	integer	az64	[]	0
9	user_dilab_student5	fct_sales	date_id	date	az64	[]	0
10	user_dilab_student5	fct_sales	unit_cost	integer	az64	[]	0
11	user_dilab_student5	fct_sales	unit_price	integer	az64	[]	0
12	user_dilab_student5	fct_sales	sales_quantity	integer	az64	[]	0

```
select "schema", "table", diststyle
from SVV_TABLE_INFO
where "schema"='user_dilab_student5'
and "table" ='fct_sales';
```

Results 1 x

select "schema", "table", diststyle from SVV_TABLE_INFO where "schema"

schema	table	diststyle
user_dilab_student5	fct_sales	AUTO(EVEN)

As seen before, redshift didn't provide any sort keys or dist keys.
distribution style for large tables is "EVEN", while small tables have "ALL"

A. Identify compression types I
decided to take dim_customers.

```
select pg.schemaname,
pg.tablename,
pg.column,
pg.type,
pg.encoding
from pg_table_def pg
where schemaname='user_dilab_student5'
and tablename = 'dim_customers';
```

pg_table_def 1 X

select pg.schemaname, pg.tablename, pg.column, pg.type, pg.encoding from

	asc schemaname	asc tablename	asc column	type	asc encoding
1	user_dilab_student5	dim_customers	customer_surr_id	integer	az64
2	user_dilab_student5	dim_customers	source_system	character varying(50)	lzo
3	user_dilab_student5	dim_customers	source_entity	character varying(50)	lzo
4	user_dilab_student5	dim_customers	customer_src_id	character varying(50)	lzo
5	user_dilab_student5	dim_customers	customer_first_name	character varying(50)	lzo
6	user_dilab_student5	dim_customers	customer_last_name	character varying(50)	lzo
7	user_dilab_student5	dim_customers	customer_gender	character varying(50)	lzo
8	user_dilab_student5	dim_customers	customer_birthday	date	az64
9	user_dilab_student5	dim_customers	customer_phone	character varying(50)	lzo
10	user_dilab_student5	dim_customers	update_dt	date	az64
11	user_dilab_student5	dim_customers	insert_dt	date	az64

Checking data in no_comp table

```
select *
from user_dilab_student5.dim_customers_no_comp
limit 20;
```

_customers_no_comp 1 X

select * from user_dilab_student5.dim_customers_no_comp limit 20

customer_surr_id	source_system	source_entity	customer_src_id	customer_first_name	customer_last_name	customer_gender
192	3nf	ce_CUSTOMERS	40425	Dalia	Hamilton	Male
208	3nf	ce_CUSTOMERS	23614	Rachael	Lloyd	Female
224	3nf	ce_CUSTOMERS	64530	Nick	Weston	Male
240	3nf	ce_CUSTOMERS	82979	Rufus	Wilson	Male
256	3nf	ce_CUSTOMERS	80988	Martin	Porter	Female
272	3nf	ce_CUSTOMERS	42111	Parker	Slater	Male
288	3nf	ce_CUSTOMERS	18301	Amy	Tailor	Male
304	3nf	ce_CUSTOMERS	17763	Adalie	Cobb	Female
320	3nf	ce_CUSTOMERS	42368	Hadley	Yates	Female
336	3nf	ce_CUSTOMERS	18693	Molly	Norton	Female
352	3nf	ce_CUSTOMERS	52006	Cara	Dale	Male
368	3nf	ce_CUSTOMERS	40103	Henry	Faulkner	Female

A. Check compression types for no_comp

```
select pg.schemaname,
       pg.tablename,
       pg.column,
       pg.type,
       pg.encoding
from pg_table_def pg
where schemaname='user_dilab_student5'
and tablename = 'dim_customers_no_comp';
```

g_table_def 1 ×

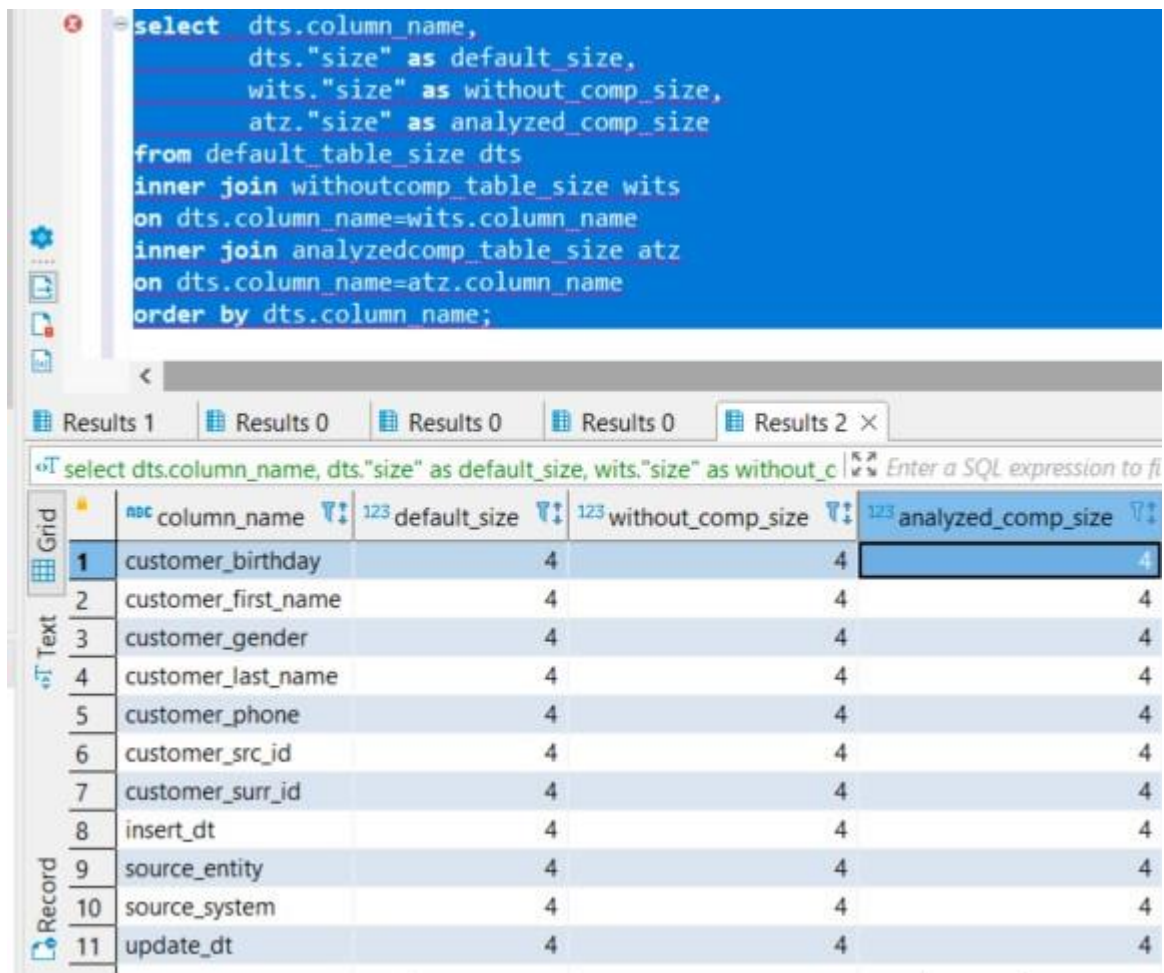
select pg.schemaname, pg.tablename, pg.column, pg.type, pg.encoding from | Enter a SQL expression to filter results (us

	asc schemaname	asc tablename	asc column	type	asc encoding
1	user_dilab_student5	dim_customers_no_comp	customer_surr_id	integer	none
2	user_dilab_student5	dim_customers_no_comp	source_system	character varying(256)	none
3	user_dilab_student5	dim_customers_no_comp	source_entity	character varying(256)	none
4	user_dilab_student5	dim_customers_no_comp	customer_src_id	character varying(256)	none
5	user_dilab_student5	dim_customers_no_comp	customer_first_name	character varying(256)	none
6	user_dilab_student5	dim_customers_no_comp	customer_last_name	character varying(256)	none
7	user_dilab_student5	dim_customers_no_comp	customer_gender	character varying(256)	none
8	user_dilab_student5	dim_customers_no_comp	customer_birthday	date	none
9	user_dilab_student5	dim_customers_no_comp	customer_phone	character varying(256)	none
10	user_dilab_student5	dim_customers_no_comp	update_dt	date	none
11	user_dilab_student5	dim_customers_no_comp	insert_dt	date	none

Use analyze command

Analyze compression user_dilab_student5.dim_customers_no_comp
comprows 87662;

B Compare size of tables with default compression, without compression and analyzed



```
select dts.column_name,
       dts."size" as default_size,
       wits."size" as without_comp_size,
       atz."size" as analyzed_comp_size
from default_table_size dts
inner join withoutcomp_table_size wits
on dts.column_name=wits.column_name
inner join analyzedcomp_table_size atz
on dts.column_name=atz.column_name
order by dts.column_name;
```

	column_name	default_size	without_comp_size	analyzed_comp_size
1	customer_birthday	4	4	4
2	customer_first_name	4	4	4
3	customer_gender	4	4	4
4	customer_last_name	4	4	4
5	customer_phone	4	4	4
6	customer_src_id	4	4	4
7	customer_surr_id	4	4	4
8	insert_dt	4	4	4
9	source_entity	4	4	4
10	source_system	4	4	4
11	update_dt	4	4	4

Redshift had chosen optimal types of compression

TASK 4. WORK WITH STORED PROCEDURE

Turning cach of

```
ALTER USER dilab_student5 SET enable_result_cache_for_session TO off;
commit;
```

```
select "schema",
       "table",
       diststyle,
       sortkey1
from SVV_TABLE_INFO
where "schema"='user_dilab_student5'
and ("table"='fct_sales'
or "table"='dim_dates'
or "table"='dim_payment_types'
or "table"='dim_stores'
or "table"='dim_customers'
or "table"='dim_employees');
```

Results 1	Results 0	Results 0	Results 0	Results 2 ×
ct "schema", "table", diststyle, sortkey1 from SVV_TABLE_INFO where Enter c				
noc schema	noc table	noc diststyle	noc sortkey1	
user_dilab_student5	dim_dates	AUTO(ALL)	AUTO(SORTKEY)	
user_dilab_student5	dim_employees	AUTO(ALL)	AUTO(SORTKEY)	
user_dilab_student5	fct_sales	AUTO(EVEN)	AUTO(SORTKEY)	
user_dilab_student5	dim_customers	AUTO(EVEN)	AUTO(SORTKEY)	
user_dilab_student5	dim_stores	AUTO(ALL)	AUTO(SORTKEY)	
user_dilab_student5	dim_payment_types	AUTO(ALL)	AUTO(SORTKEY)	

Then i altered tables to set sortkeys

```
-----alter diststyle and sort keys-----

alter table user_dilab_student5.fct_sales alter sortkey none;
alter table user_dilab_student5.dim_dates alter sortkey none;
alter table user_dilab_student5.dim_stores alter sortkey none;
alter table user_dilab_student5.dim_customers alter sortkey none;
alter table user_dilab_student5.dim_employees alter sortkey none;
alter table user_dilab_student5.dim_payment_types alter sortkey none;

alter table user_dilab_student5.dim_payment_types alter diststyle even;
alter table user_dilab_student5.fct_sales alter diststyle even;
alter table user_dilab_student5.dim_dates alter diststyle all;
alter table user_dilab_student5.dim_stores alter diststyle all;
alter table user_dilab_student5.dim_customers alter diststyle even;
alter table user_dilab_student5.dim_employees alter diststyle even;

commit;
```



```
select "schema",
       "table",
       diststyle,
       sortkey1
from SVV_TABLE_INFO
where "schema"='user_dilab_student5'
and ("table"='fct_sales'
or "table"='dim_dates'
or "table"='dim_payment_types'
or "table"='dim_stores'
or "table"='dim_customers'
or "table"='dim_employees');
```

```

-- explain
select distinct ds.store_name, f.date_id,
dpt.payment_type,
sum (f.sales_quantity) as quantity_sold,
sum (f.unit_price) as price_per_unit,
sum (f.unit_price* f.sales_quantity) as total_revenue
from user_dilab_student5.fct_sales f
left join user_dilab_student5.dim_payment_types dpt on dpt.payment_type_surr_id = f.payment_type_surr_id
left join user_dilab_student5.dim_stores ds on f.store_surr_id = ds.store_surr_id
where date_id >= '2022-03-17'
group by f.date_id, f.product_surr_id, dpt.payment_type, ds.store_name
order by total_revenue desc
limit 20;

```

Save Cancel Sprint 50 20 Rows: 1 20 row(s) fetched

ABC QUERY PLAN	
XN Merge (cost=1000000497963.09..1000000498248.18 rows=114036 width=33)	
Merge Key: sum((f.sales_quantity * f.unit_price))	
-> XN Network (cost=1000000497963.09..1000000498248.18 rows=114036 width=33)	
Send to leader	
-> XN Sort (cost=1000000497963.09..1000000498248.18 rows=114036 width=33)	
Sort Key: sum((f.sales_quantity * f.unit_price))	
-> XN Unique (cost=485533.66..488384.56 rows=114036 width=33)	
-> XN HashAggregate (cost=485533.66..486674.02 rows=114036 width=33)	
-> XN Hash Left Join DS_DIST_ALL_NONE (cost=0.66..467687.20 rows=1019798 width=33)	
Hash Cond: ("outer".store_surr_id = "inner".store_surr_id)	
-> XN Hash Left Join DS_BCAST_INNER (cost=0.03..434993.06 rows=999801 width=27)	
Hash Cond: ("outer".payment_type_surr_id = "inner".payment_type_surr_id)	
-> XN Seq Scan on fct_sales f (cost=0.00..12497.51 rows=999801 width=24)	
Filter: (date_id >= '2022-03-17'::date)	
-> XN Hash (cost=0.02..0.02 rows=2 width=11)	
-> XN Seq Scan on dim_payment_types dpt (cost=0.00..0.02 rows=2 width=11)	
-> XN Hash (cost=0.51..0.51 rows=51 width=14)	
-> XN Seq Scan on dim_stores ds (cost=0.00..0.51 rows=51 width=14)	

Runs	Time	Cost
1	43	1000000497963.09..1000000498248.18
2	31	1000000497963.09..1000000498248.18
3	32	1000000497963.09..1000000498248.18

TASK 5. WORK WITH OPTIMIZATION OF DISTRIBUTION STYLE AND SORT KEYS.

After that I decided to add sorkeys to fct_sales and dim stores as seen below.

```
alter table user_dilab_student5.fct_sales alter sortkey (unit_price,sales_quantity);
alter table user_dilab_student5.dim_stores alter sortkey (store_name);
```

Runs	Time	Cost
1	40	1000000497963.09..1000000497223.22
2	29	1000000497963.09..1000000497223.22
3	27	1000000497963.09..1000000497223.22

Extremely small improvement can be seen.

TASK 6. RUN THE STORED PROCEDURE USING OPTIMIZED TABLES AND LOAD DATA TO YOUR REPORT

```

as $$
declare
integer_val integer:=0;
begin
raise info 'Procedure update_report started';
insert into user_dilab_student5.report (store_name,
                                         date_id,
                                         payment_type,
                                         quantity_sold,
                                         price_per_unit,
                                         total_revenue)

select distinct ds.store_name,
               f.date_id,
               dpt.payment_type,
               sum (f.sales_quantity) as quantity_sold,
               sum (f.unit_price) as price_per_unit,
               sum (f.unit_price* f.sales_quantity) as total_revenue
from user_dilab_student5.fct_sales f
left join user_dilab_student5.dim_payment_types dpt
on dpt.payment_type_surr_id = f.payment_type_surr_id
left join user_dilab_student5.dim_stores ds
on f.store_surr_id = ds.store_surr_id
where date_id >= '2022-03-17'
group by      f.date_id,
               f.product_surr_id,
               dpt.payment_type,
               ds.store_name

order by total_revenue desc
get diagnostics integer_val := row_count;
raise notice 'Table report inserted, rows: %',integer_val ;
exception
when others then raise exception 'error message: %, error code: %',sqlerrm, sqlstate;
end;
$$ language plpgsql;

select *
from user_dilab_student5.report;

```

report 1 ×

select * from user_dilab_student5.report Enter a SQL expression to filter results (use Ctrl+Space)

	asc store_name	date_id	asc payment_type	123 quantity_sold	123 price_per_unit	123 total_revenue
1	Nime	2022-03-17	CASH	5,672	550	28,360
2	Man Mobile	2022-03-17	CASH	5,339	525	26,695
3	ReadABook	2022-03-17	CASH	5,148	500	25,740
4	Ingocal	2022-03-17	CASH	4,904	475	24,520
5	Caint	2022-03-17	CASH	4,804	475	24,020
6	MemoMe	2022-03-17	CASH	4,683	450	23,415
7	Paize	2022-03-17	CASH	4,667	450	23,335
8	BeReady	2022-03-17	CASH	4,625	425	23,125
9	Koob	2022-03-17	CASH	4,579	425	22,895

COPY QUESTION

Created tables lineorder_1 and lineorder_2

```
CREATE TABLE lineorder_1
(
  lo_orderkey bigint NOT null ,
  lo_linenumber bigint NOT NULL,
  lo_custkey bigint NOT NULL,
  lo_partkey bigint NOT NULL,
  lo_suppkey bigint NOT NULL,
  lo_orderdate bigint NOT NULL,
  lo_orderpriority VARCHAR(20) NOT NULL,
  lo_shippriority VARCHAR(5) NOT NULL,
  lo_quantity bigint NOT NULL,
  lo_extendedprice bigint NOT NULL,
  lo_ordertotalprice bigint NOT NULL,
  lo_discount bigint NOT NULL,
  lo_revenue bigint NOT NULL,
  lo_supplycost bigint NOT NULL,
  lo_tax bigint NOT NULL,
  lo_commitdate bigint NOT NULL,
  lo_shipmode VARCHAR(15) NOT NULL
);

CREATE TABLE lineorder_2
(
  lo_orderkey bigint NOT null ,
  lo_linenumber bigint NOT NULL,
  lo_custkey bigint NOT NULL,
  lo_partkey bigint NOT NULL,
  lo_suppkey bigint NOT NULL,
  lo_orderdate bigint NOT NULL,
  lo_orderpriority VARCHAR(20) NOT NULL,
  lo_shippriority VARCHAR(5) NOT NULL,
  lo_quantity bigint NOT NULL,
  lo_extendedprice bigint NOT NULL,
  lo_ordertotalprice bigint NOT NULL,
  lo_discount bigint NOT NULL,
  lo_revenue bigint NOT NULL,
  lo_supplycost bigint NOT NULL,
  lo_tax bigint NOT NULL,
  lo_commitdate bigint NOT NULL,
  lo_shipmode VARCHAR(15) NOT NULL
);
```

2. Loading data to tables

```

copy lineorder_1
from 's3://dilabbucket/files/lineorder_file/'
credentials
'aws_iam_role=arn:aws:iam::260586643565:role/dilab-redshift-role'
region 'eu-central-1'
delimiter ','
gzip
DATEFORMAT AS 'auto'
IGNOREHEADER 1;

select * from stl_load_errors;

copy lineorder_2
from 's3://dilabbucket/files/lineorders/'
credentials
'aws_iam_role=arn:aws:iam::260586643565:role/dilab-redshift-role'
format as parquet;

insert into temp_table_times (querytxt, starttime, endtime,diff)
select querytxt, starttime, endtime,datediff(millisecond,starttime,endtime)
from stl_query where userid=109
and querytxt like 'copy lineorder_1%'
order by starttime desc limit 20;
insert into temp_table_times (querytxt, starttime, endtime,diff)
select querytxt, starttime, endtime,datediff(millisecond,starttime,endtime)
from stl_query where userid=109
and querytxt like 'copy lineorder_2%'
order by starttime desc limit 20;

```

querytxt	starttime	endtime	diff
copy lineorder_1 from 's3://dilabbucket/files/lineorder_file/' credentials "" region 'eu-central-1' delimiter ',' gzip DATEFORMAT	2022-04-03 10:30:40.091	2022-04-03 10:35:14.899	274,808
copy lineorder_2 from 's3://dilabbucket/files/lineorders/' credentials "" format as parquet	2022-04-03 10:27:03.537	2022-04-03 10:29:03.220	119,683

Type of file	Number of files	Size of files	Load time, s	Count of rows
gz	1	3 Gb	274,808	100.000.008
parquet	4	2.22 Gb	119,683	100.000.008

The reason why the first query works longer is that it reads compressed data. But i was surprised that while file was compressed by 25%, the actual read time is 2 times longer.

EXTERNAL TABLES

1.

```
-----create external schema-----
create external schema if not exists user_dilab_student5_ext
from data catalog
database 'artiom_dolzhenko_database'
IAM_ROLE 'arn:aws:iam::260586643565:role/dilab-redshift-role'
commit;

-----check external schema-----
select * from pg_catalog.svv_external_schemas
where schemaname = 'user_dilab_student5_ext';

-----upload data to s3-----
unload ('Select * from fct_sales
        Where extract (year from date_id)=2022
        and extract (MONTH from date_id)=1;')
to 's3://artiom-dolzhenko-aws-bucket/stores/bl_dm/sales_2022-01/'
iam_role 'arn:aws:iam::260586643565:role/dilab-redshift-role';
unload ('Select * from fct_sales
        Where extract (year from date_id)=2022
        and extract (MONTH from date_id)=2;')
to 's3://artiom-dolzhenko-aws-bucket/stores/bl_dm/sales_2022-02/'
iam_role 'arn:aws:iam::260586643565:role/dilab-redshift-role';
unload ('Select * from fct_sales
        Where extract (year from date_id)=2022
        and extract (MONTH from date_id)=3;')
to 's3://artiom-dolzhenko-aws-bucket/stores/bl_dm/sales_2022-03/'
iam_role 'arn:aws:iam::260586643565:role/dilab-redshift-role';

-----create external table-----
CREATE external TABLE user_dilab_student5_ext.FCT_SALES_EXT_PART
(
  source_system      VARCHAR(50) ,
  source_entity      VARCHAR(50) ,
  product_surr_id    INTEGER ,
  payment_type_surr_id INTEGER ,
  promotion_surr_id  INTEGER ,
  store_surr_id      INTEGER ,
  employee_surr_id   INTEGER ,
  customer_surr_id   INTEGER ,
  date_id            DATE ,
  unit_cost          INTEGER ,
  unit_price         INTEGER ,
  sales_quantity     INTEGER ,
  update_dt          DATE ,
  insert_dt          DATE
)
partitioned by (DATE_ID DATE)
row format delimited
fields terminated by '|'
stored as textfile
location 's3://artiom-dolzhenko-aws-bucket/stores/bl_dm/fct_sales/'
table properties ('numRows'='999801');
```